

Comprehensive ASR Prompt Development Report

Evaluating Multimodal LLMs on Code-Switched Nepali Audio

August 14, 2026

Abstract

This research report details the systematic prompt development and evaluation of four multimodal Large Language Models (LLMs) for Automatic Speech Recognition (ASR). The primary objective is to transcribe code-switched speech (mixing Nepali and English) verbatim, adhering to strict linguistic rules: transcribing Nepali speech in Devanagari script and English speech in Latin script. The evaluation was conducted across two discrete stages to prevent prompt overfitting and to finalize the optimal model and prompt configuration for a massive 500-utterance final benchmark.

1 Introduction and Methodology

1.1 Research Objective

The goal of this study is to evaluate the zero-shot and few-shot capabilities of multimodal LLMs when tasked with complex audio transcription. Unlike traditional ASR, this task demands rigorous formatting adherence. Specifically, the model must maintain strict orthographic separation based on spoken language (Devanagari for Nepali, Latin for English) while ignoring the urge to grammatically correct or summarize the transcriptions.

1.2 Testing Stages

To ensure the chosen prompt generalizes to unseen audio and does not suffer from prompt-overfitting, the dataset was split into two stages:

- **Stage 1 (Pilot Testing):** 30 audio files representing clean (10), code-switched (10), and noisy (10) audio conditions. Used to test a wide variety of prompts and establish baseline performance.
- **Stage 2 (Validation):** 60 unseen audio files with similar distributions. Used to validate the winning models from Stage 1.

1.3 Prompt Complexity Levels

Each model was tested using 5 distinct prompting strategies, progressively increasing in guidance:

- **L0 (Minimal Control):** Zero-shot. No formatting rules provided; basic transcription request.
- **L1 (Plain Instruction):** Zero-shot. Standard transcription instructions with minor constraints.
- **L2 (Linguistic Rules):** Zero-shot. Highly detailed formatting rules explicitly forbidding summarization and demanding script separation.
- **L3 (1-Shot):** Detailed rules paired with exactly 1 perfect text exemplar.
- **L4 (3-Shot):** Detailed rules paired with 3 perfect text exemplars.

1.4 Metrics

- **Word Error Rate (WER) & Character Error Rate (CER):** Standard ASR metrics. Lower scores indicate better transcription accuracy.
- **Script Match Rate:** The percentage of words correctly formatted in their target alphabet. Higher is better.
- **Hallucination Rate:** The frequency of the model generating fabricated text not present in the audio. Lower is better.

2 Stage 1: Pilot Testing Findings

During Stage 1, all four models (Voxtral-Mini-3B, Gemma4-12B, Qwen2.5-Omni-3B, and Phi-4-Multimodal) were evaluated to weed out underperformers.

Table 1: Stage 1 Pilot Test: Best Configuration per Model

Rank	Model Candi- date	Optimal Setup	WER	Script Match	Hallucination Rate
#1	Voxtral-Mini-3B	L3 (1-Shot)	0.52	74.0%	0.0%
#2	Gemma4-12B	L2 (Zero-Shot)	0.74	77.7%	0.0%
#3	Qwen2.5-Omni-3B	L1 (Zero-Shot)	0.87	66.7%	0.0%
#4	Phi-4 Multi-modal	L3 (1-Shot)	1.28	66.7%	14.8%

2.1 Model-Specific Analysis (Stage 1)

- **Voxtral-Mini-3B:** Even its worst prompt (L1, WER 0.62) outperformed the best prompts of all other candidates. It perfectly leveraged few-shot examples (L3/L4) to drop its WER to 0.52 and boost formatting adherence, all while maintaining a 0% hallucination rate.
- **Gemma4-12B:** Exhibited the best zero-shot formatting adherence across all models. At L2 (detailed rules), it achieved a 77.7% script match and a 0.74 WER. However, introducing an exemplar (L3) caused catastrophic failure, spiking its WER to 1.30 and triggering severe hallucinations (25.9%).
- **Qwen2.5-Omni-3B:** Performed best with simple instructions (L1, WER 0.87). Like Gemma4, complex rules and few-shot examples degraded its accuracy.
- **Phi-4 Multimodal:** Required the most support. It completely failed to format text zero-shot (L0: 0% script match). Even with a 1-shot example (L3), it peaked at a poor 1.28 WER and hallucinated frequently.

3 Stage 2: Validation on Unseen Data

To test for overfitting, the models were evaluated against the 60 unseen utterances of Stage 2.

Table 2: Stage 2 Validation Test: Best Configuration per Model

Rank	Model Candi- date	Optimal Setup	WER	Script Match	Hallucination Rate
#1	Voxtral-Mini-3B	L4 (3-Shot)	0.55	68.4%	0.0%
#2	Gemma4-12B	L1 (Zero-Shot)	0.74	68.4%	1.8%
#3	Qwen2.5-Omni-3B	L2 (Zero-Shot)	1.21	63.2%	3.5%
#4	Phi-4 Multi-modal	L0 (Minimal)	1.54	0.0%	7.0%

3.1 Degradation Analysis (Stage 2)

- **Voxtral’s Consistency:** Voxtral proved highly resistant to overfitting, maintaining a 0.55 WER on the larger dataset. It continued to leverage few-shot examples flawlessly to maintain a perfect 0.0% hallucination rate.
- **Gemma4’s Rigidity:** Gemma4 maintained its excellent zero-shot capability, holding a 0.74 WER at L1. It remains the premier choice if a zero-shot pipeline is strictly mandated.
- **Qwen2.5’s Degradation:** The larger validation set exposed Qwen’s lack of generalization. Its optimal WER degraded drastically from 0.87 (Stage 1) to 1.21 (Stage 2), indicating heavy overfitting during the pilot test.
- **Phi-4’s Collapse:** Phi-4’s performance degraded entirely. Its lowest error rate (1.54) was achieved only by completely disregarding all formatting rules (0% script match).

4 Deep Comparative Analysis

4.1 Zero-Shot vs. Few-Shot Capabilities

The dataset revealed a stark divide in how modern LLMs handle context:

- **The Flexible Learners (Voxtral):** Voxtral successfully parses both explicit rules and implicit patterns from exemplars. Providing it with 3 examples (L4) consistently reduced its Word Error Rate and stabilized its output across both stages.

- **The Rigid Rule-Followers (Gemma4 & Qwen2.5):** These models excel at parsing structured linguistic rules zero-shot (L2). However, they suffer from “exemplar confusion.” When shown an example of a transcription, they attempt to map the new audio directly to the content of the exemplar, causing massive degradation in WER.

4.2 Hallucination Resistance

Hallucinations (fabricating text) are fatal in ASR.

- **Voxtral-Mini-3B** was the only model to maintain a flawless 0.0% hallucination rate across its optimal prompts in both stages.
- **Gemma4-12B** was highly hallucination-resistant in zero-shot environments, but fabricated wildly (up to 25.9% of the time) when exposed to few-shot prompts.
- **Phi-4** proved inherently unsafe for this task, frequently hallucinating regardless of prompt structure.

5 Final Verdict and Conclusion

Based on the comprehensive two-stage evaluation, the clear and undisputed champion is **Voxtral-Mini-3B**.

Unlike the other models which suffered from exemplar-confusion or dataset overfitting, Voxtral scaled beautifully across both stages. It learns effectively from few-shot examples without degrading transcription accuracy, it completely resists hallucinations, and it reliably enforces complex code-switching formatting rules.

Actionable Recommendation: We recommend permanently “freezing” the **Level 4 (3-Shot)** prompt configuration for the **Voxtral-Mini-3B** model and proceeding immediately to the final 500-utterance evaluation benchmark.